

PySpark Performance Tuning & Optimization Guide

Production Architecture & Cluster Optimization

Author: Vivek Pandey - Senior Data Engineer

Key Topics & Production Patterns:

- Broadcast Hash Joins (BHJ) and threshold tuning
- Salting techniques for mitigating skewed partitions
- Adaptive Query Execution (AQE) in Spark 3.x
- Partitioning, Coalesce, and Repartition strategies